

# Performance Analysis of Parallel Support Vector Machines on a MapReduce Architecture

Udita Patel


Department of Computer Science

# Agenda


1. Introduction
2. Support Vector Machines
3. Data Preprocessing
4. MapReduce Programming Model
5. Parallel Algorithms
6. Results
7. Conclusions
8. Future Works
9. Acknowledgement

# Introduction


- Support Vector Machines are powerful but computational complexity increases rapidly as the number of training example increases.
- Many Parallel SVM implementations exist, but not their comparative study with any benchmark dataset.
- Mostly analyzed for binary classification.
- We use the MNIST hand written digit dataset to analyze performance and accuracy of three parallel algorithms.




# Support Vector Machines



# Increased computation : Kernels




# Increased computation : Multi Class Classification



$N$  Binary SVMs

# Increased computation : Multi Class Classification



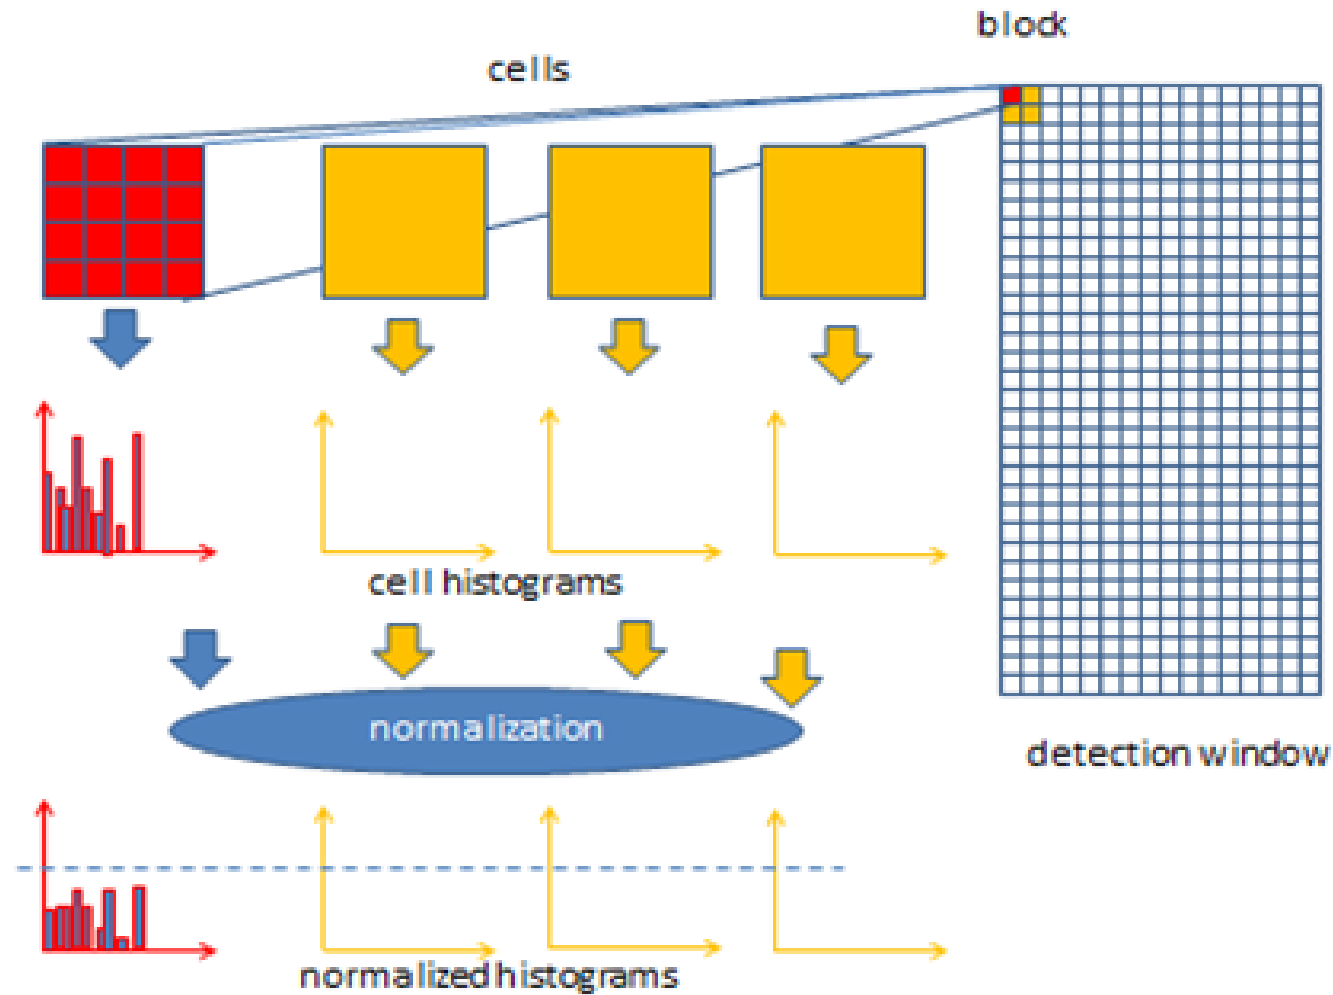
$N*(N-1)/2$  Binary SVMs

# The dataset : MNIST

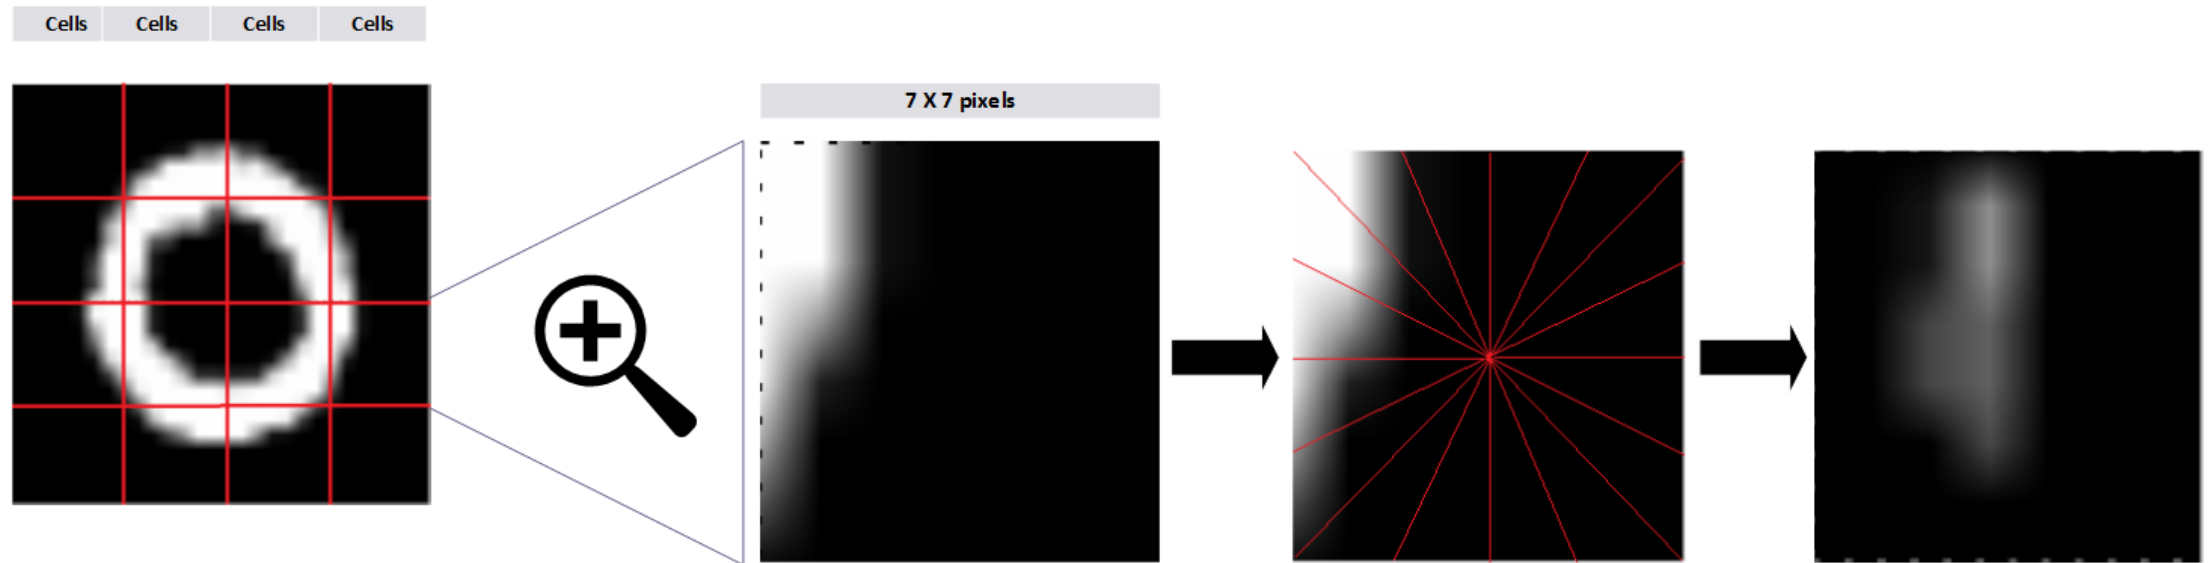
- Mixed National Institute of Standards and Technology database
- 42,000 data points
- 28 X 28 pixels



# Data Preprocessing : Histogram of Oriented Gradients (HOGs)











# Data Preprocessing : Histogram of oriented gradients




256 features

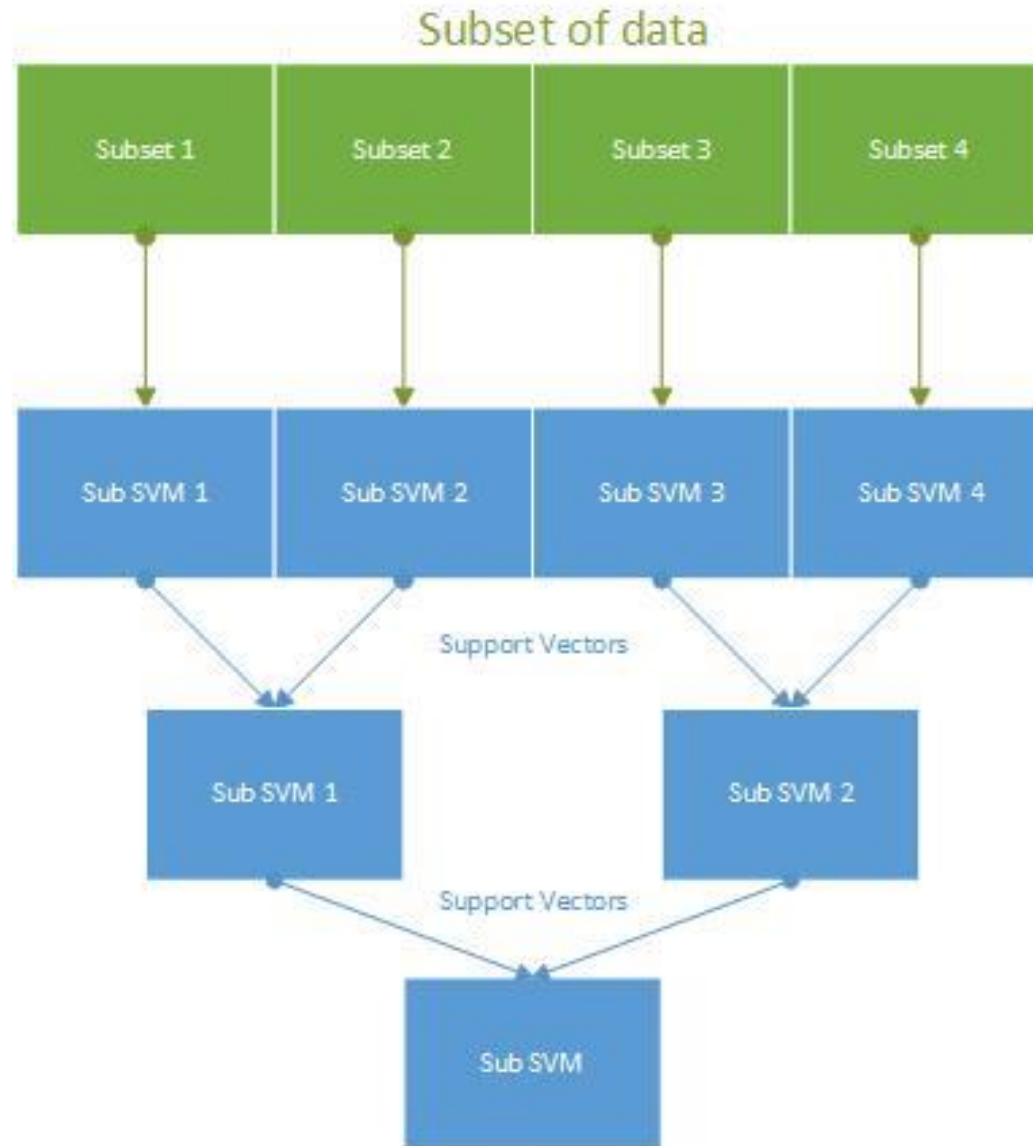
# Digits after preprocessing



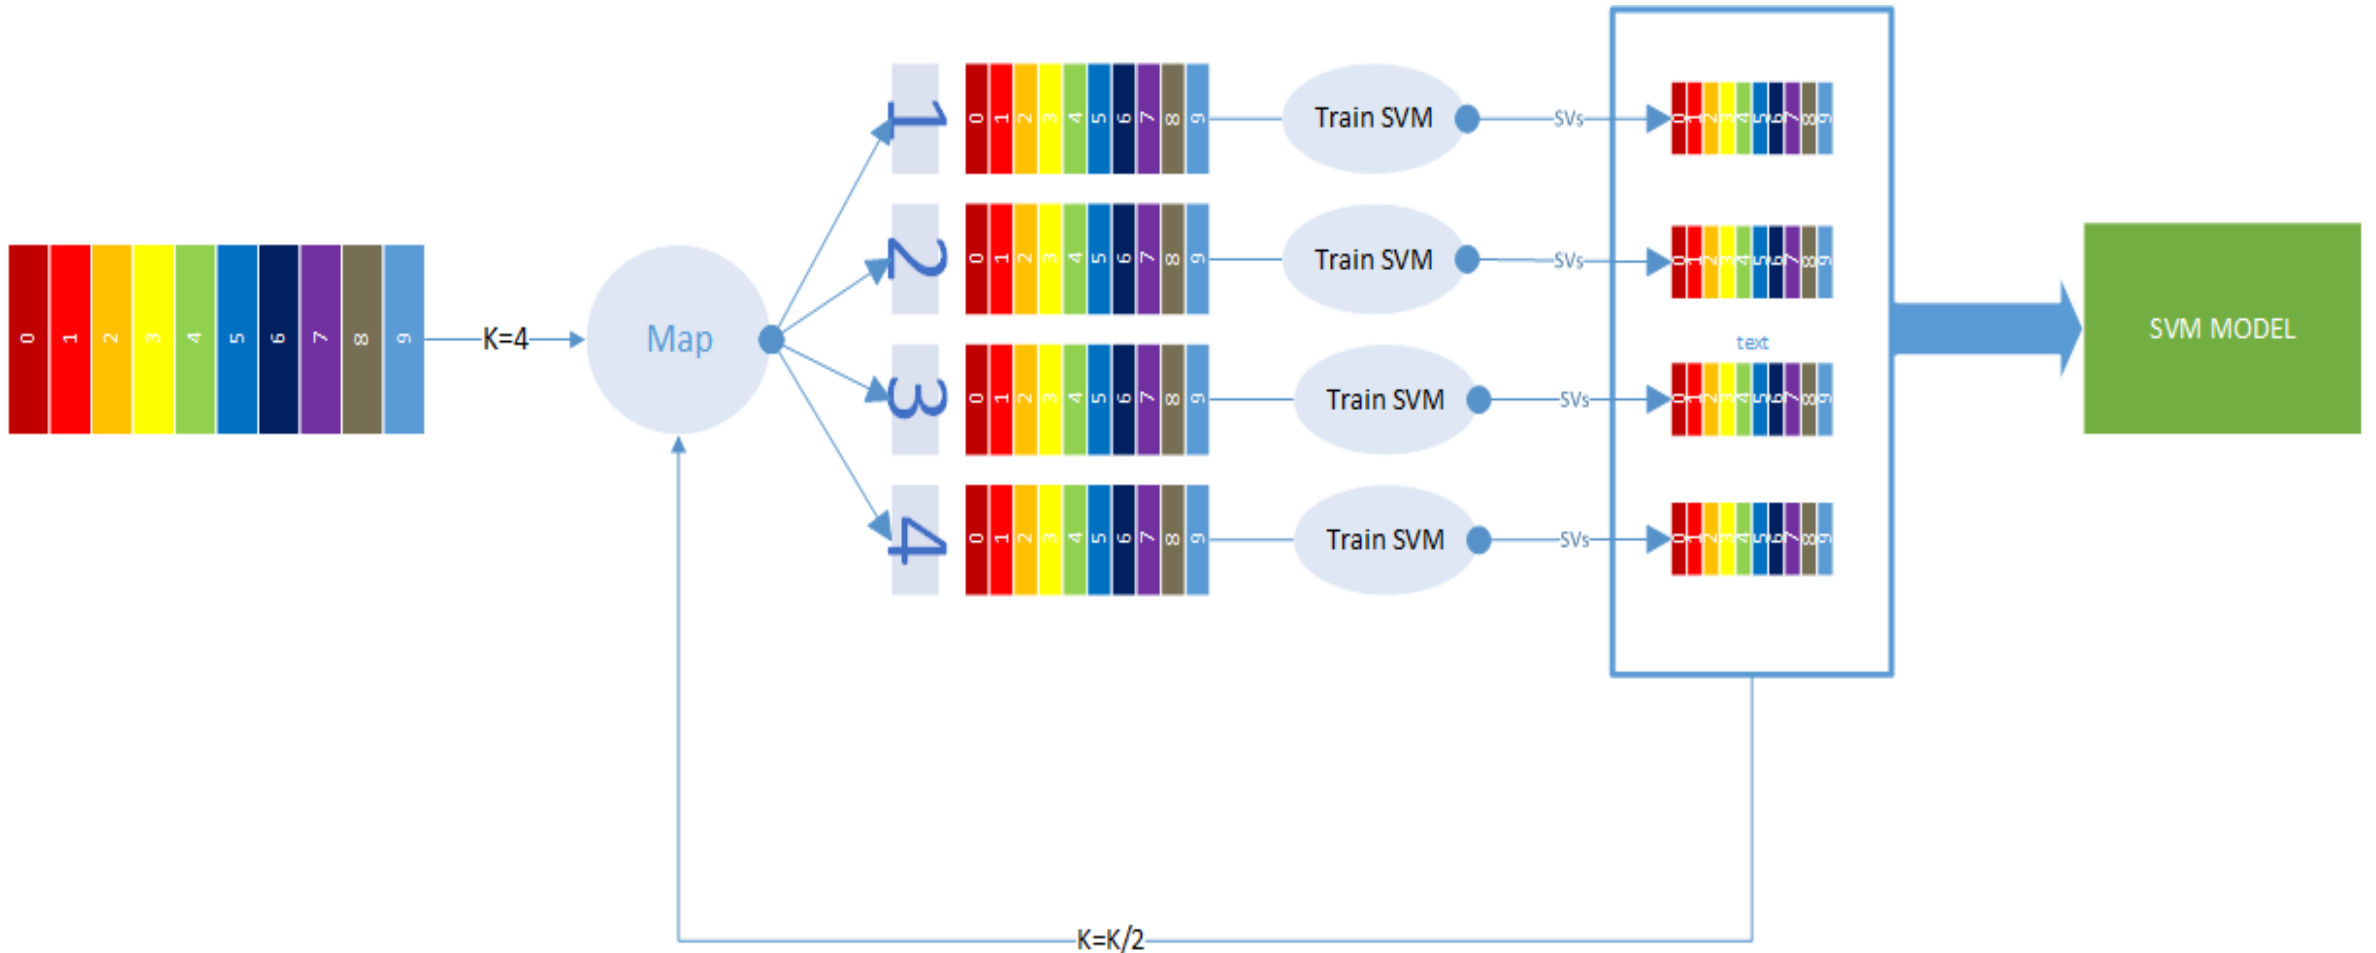
# MapReduce Model



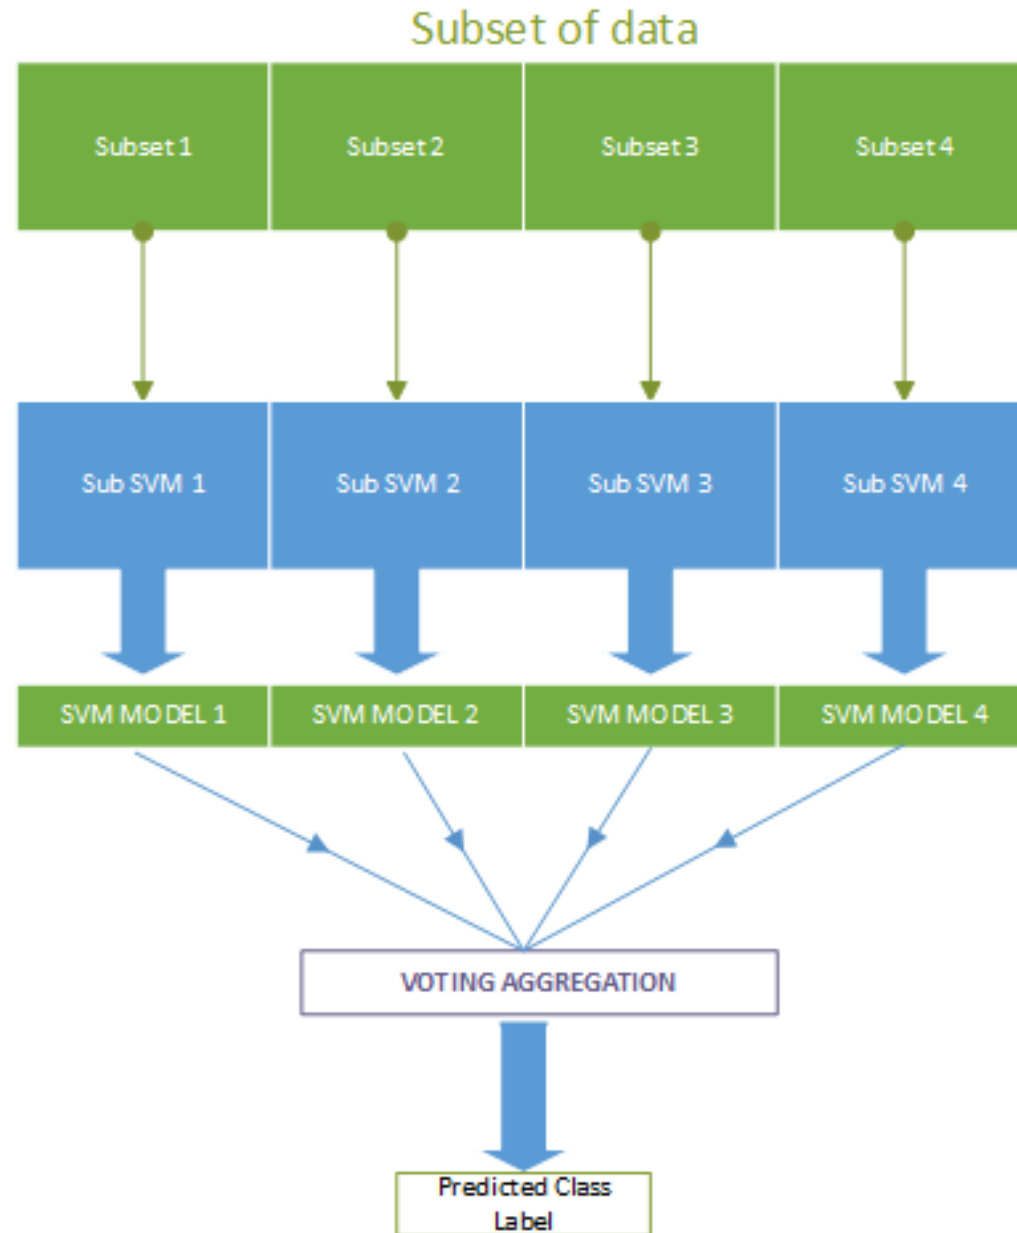
# Cascade SVM




# Cascade SVM : MapReduce Implementation



# Bagging SVM

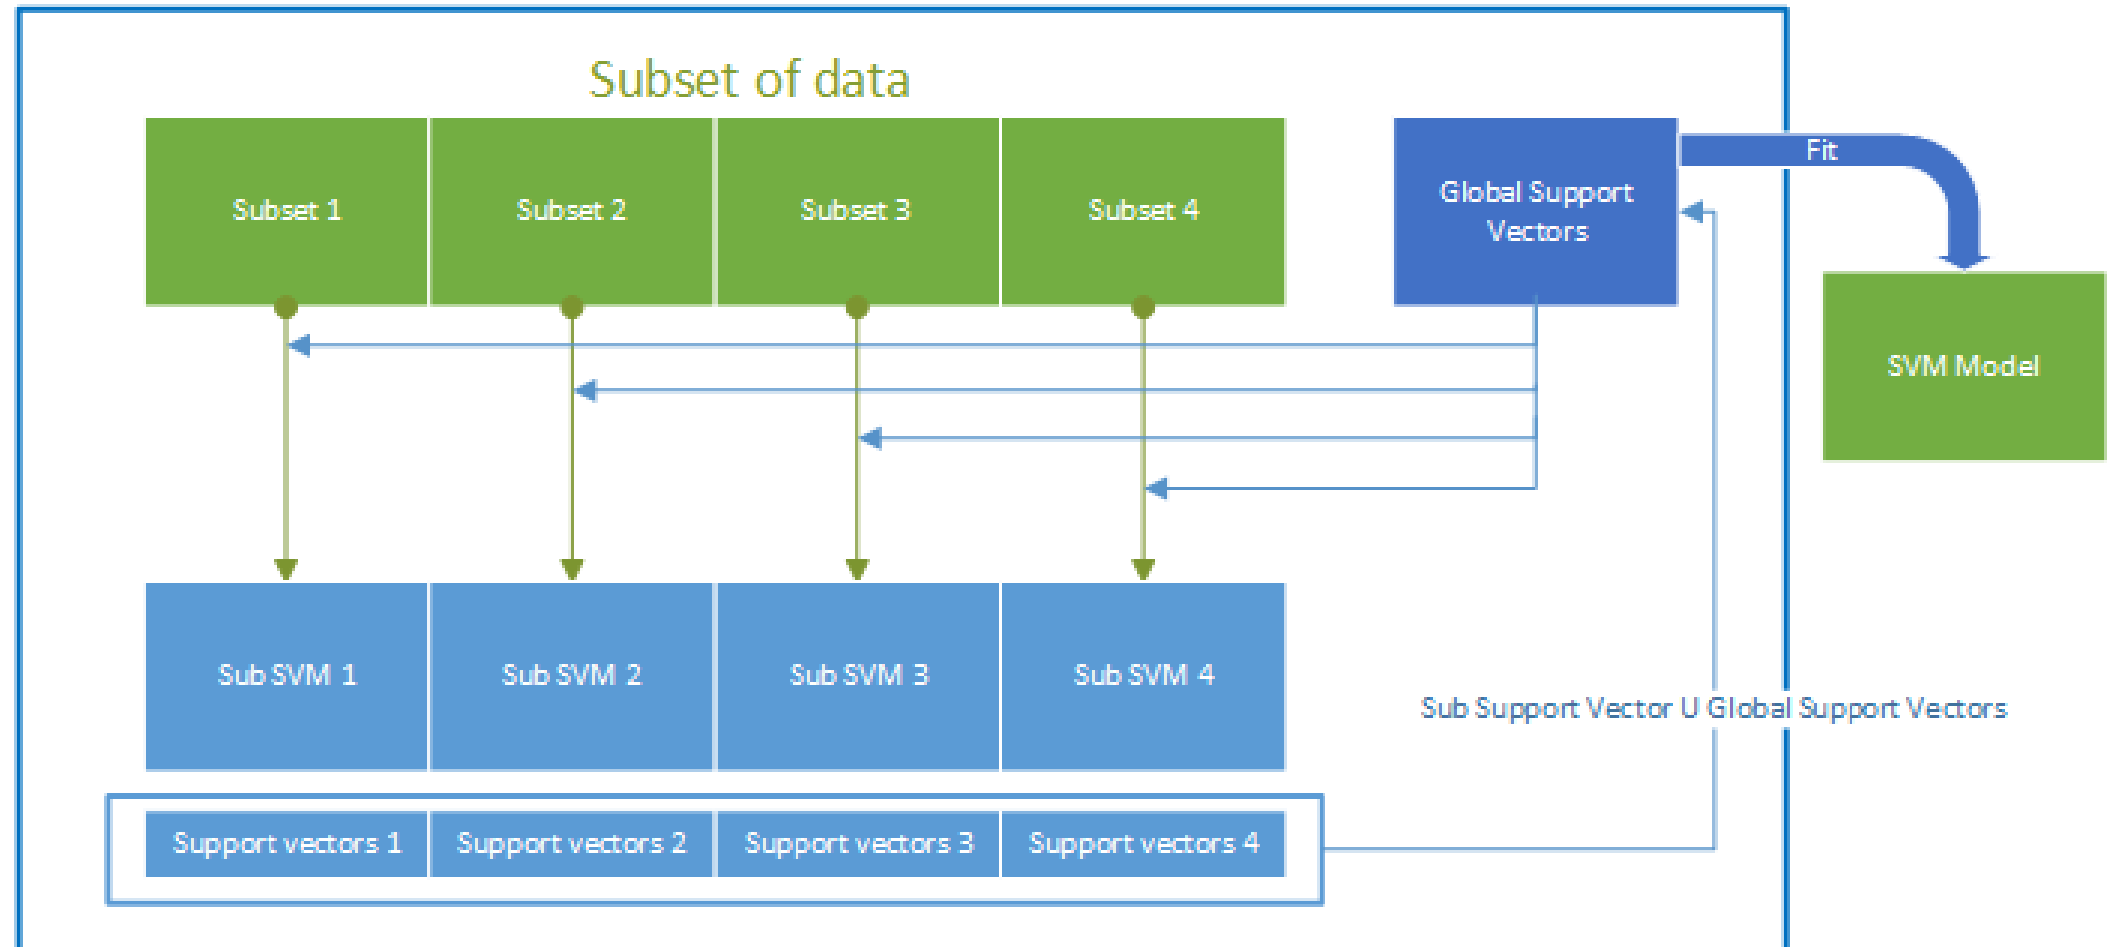


# Bagging SVM : MapReduce Implementation

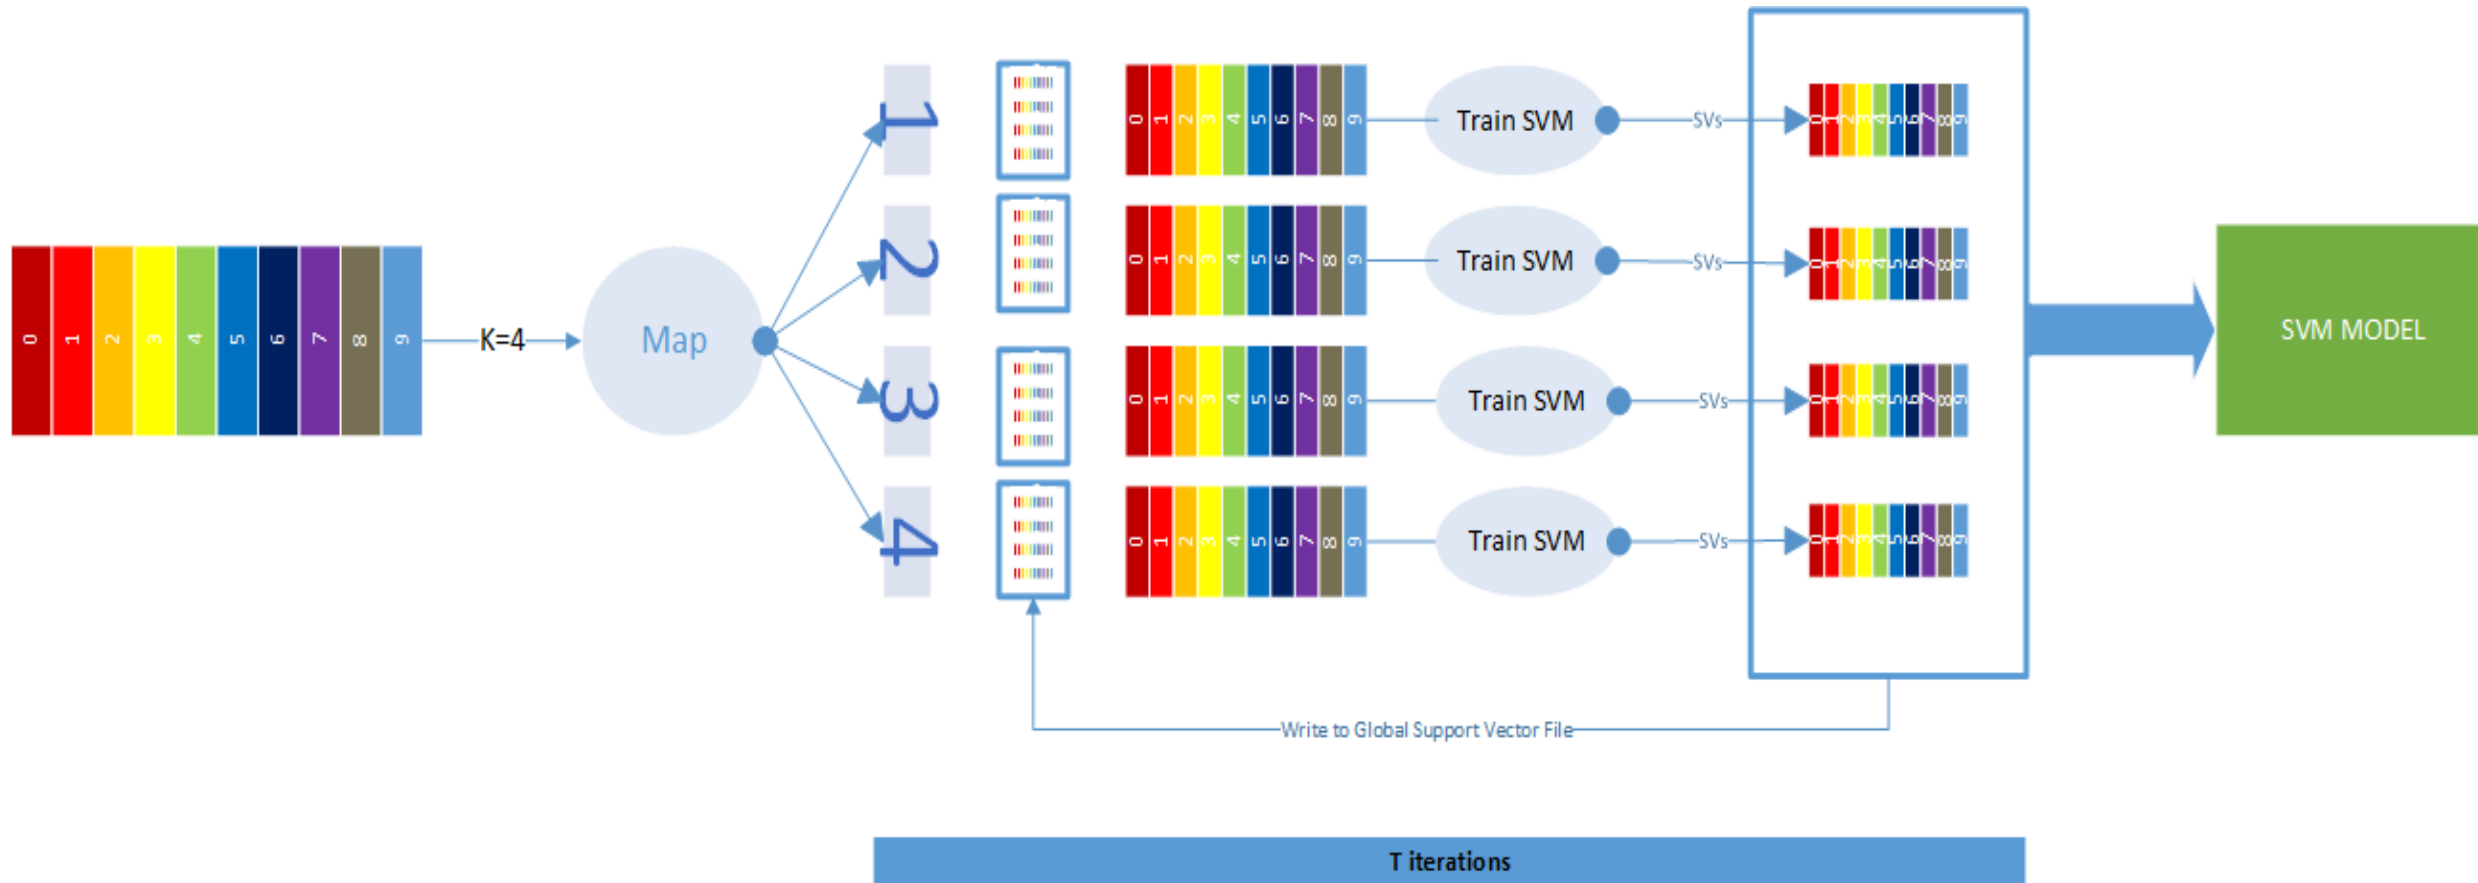





# Iterative SVM



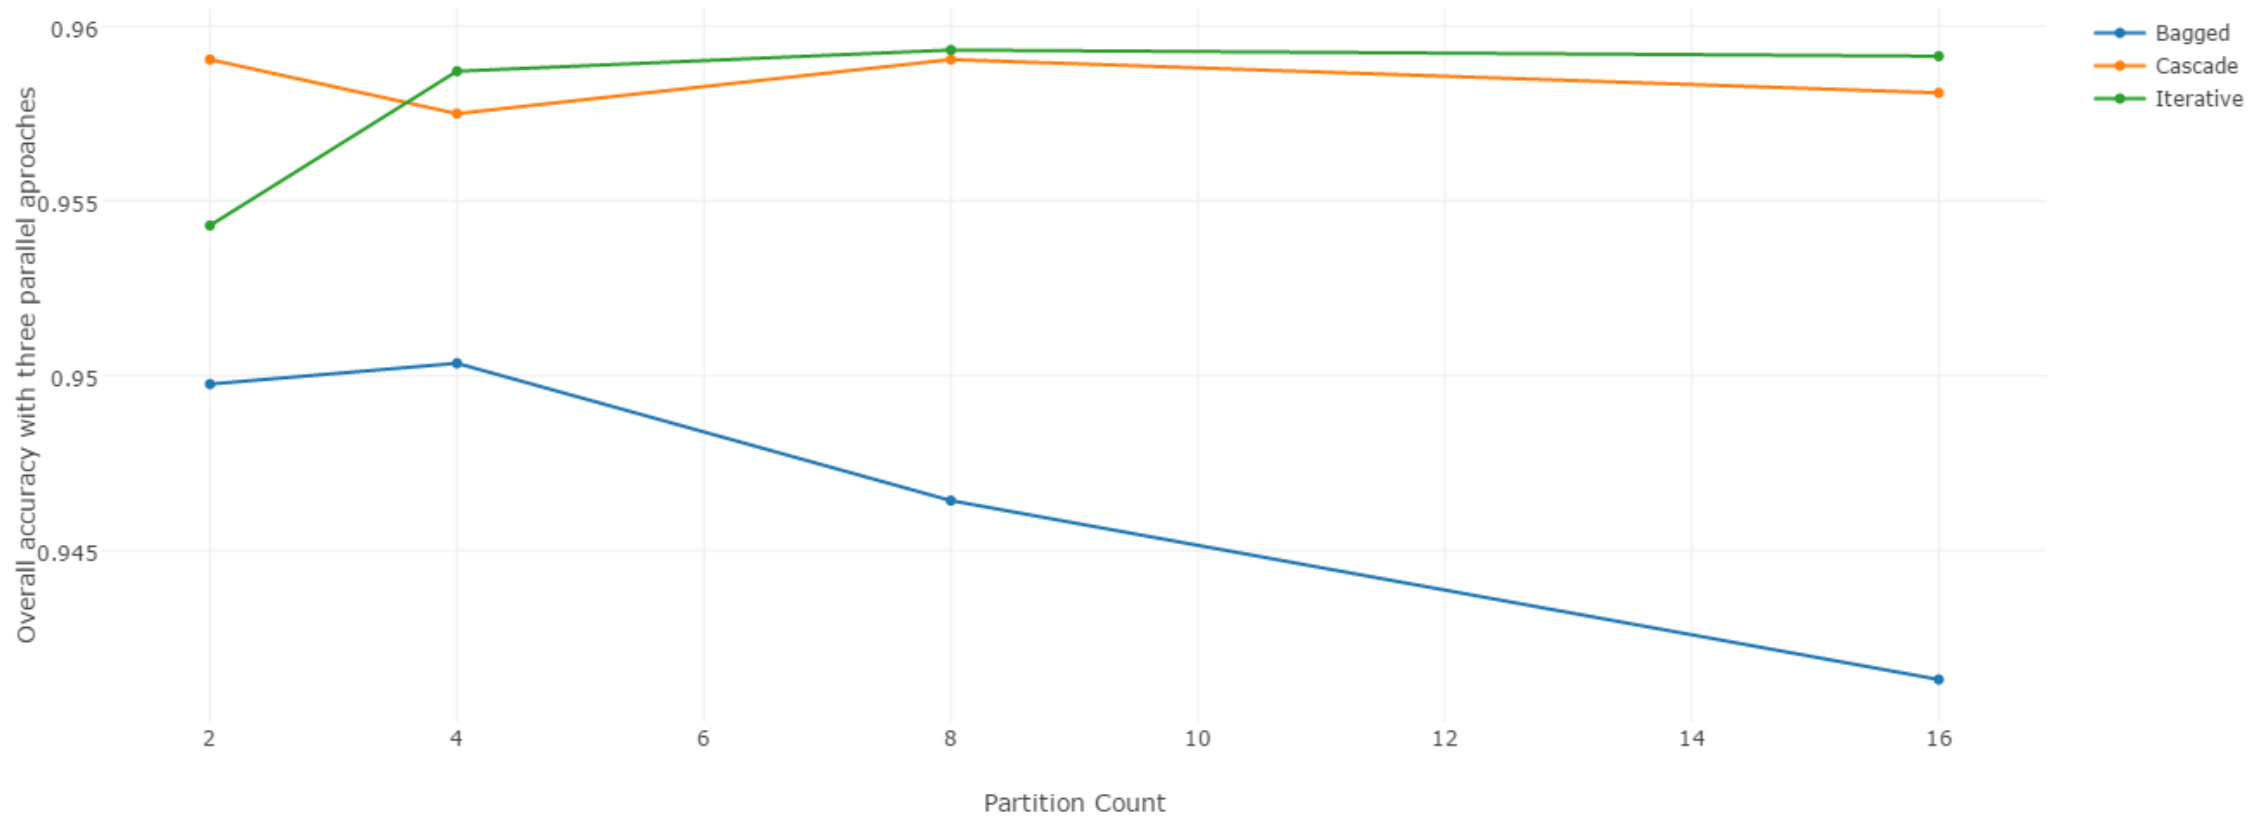
# Iterative SVM : MapReduce Implementation



# Training time vs Partition count










# Accuracy



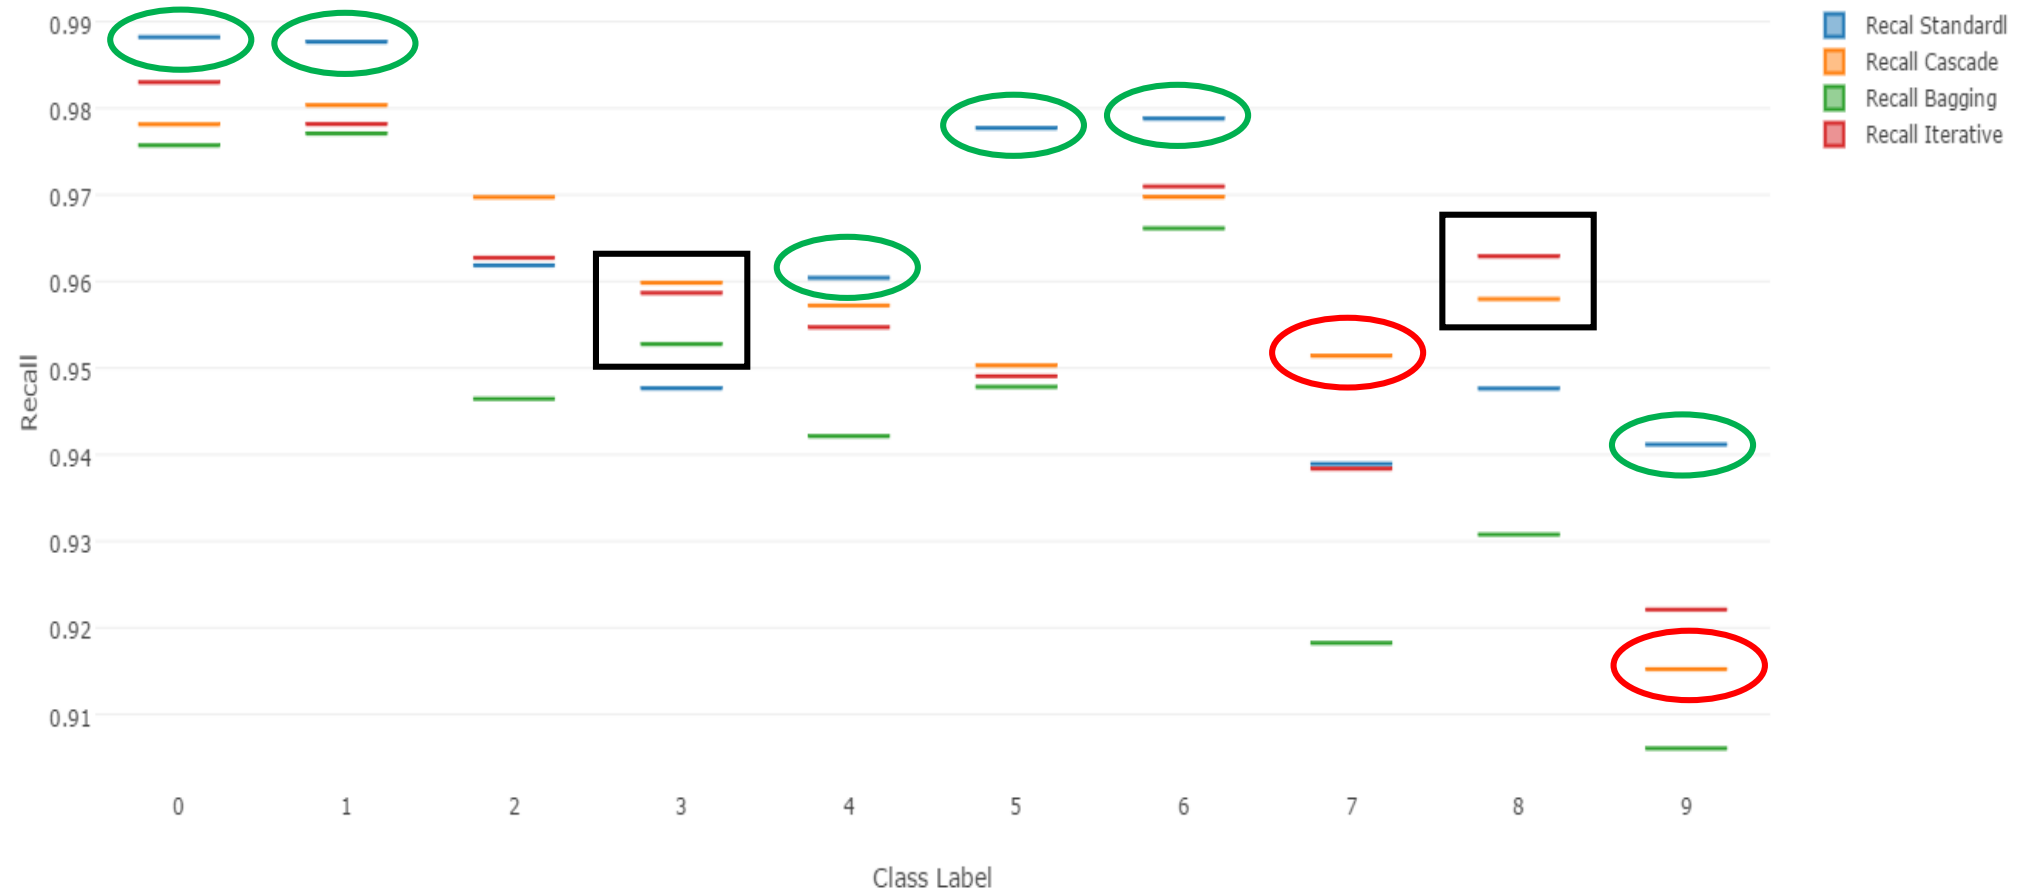
# Standard SVM confusion matrix

	0	1	2	3	4	5	6	7	8	9
0	839	2	2	0	0	0	4	0	1	1
1	1	882	1	0	2	0	1	4	1	1
2	4	4	782	9	4	5	1	2	2	0
3	4	0	11	815	1	13	0	3	10	3
4	1	1	3	0	825	0	4	4	3	18
5	0	0	2	5	1	790	3	0	6	1
6	8	1	0	0	2	4	832	0	3	0
7	0	6	4	5	7	1	0	815	3	27
8	6	8	6	2	0	8	5	1	742	5
9	1	2	1	6	15	2	1	17	3	768

# False Negatives



# Recall comparison



# Support Vectors in the final model

<b>Standard SVM</b>	<b>6340</b>
---------------------	-------------

	<b>Cascade</b>	<b>Bagging</b>	<b>Iterative</b>
2	6026	8137	7024
4	5941	10554	7748
8	5884	13909	8522
16	5888	18143	8921



# Conclusion

- The decrease in accuracy is only 0.5% to 3% for all parallel approaches, making all of them usable.
- Iterative SVM is the best choice for high partition counts.
- Bagging SVM takes lowest training time, with 3% reduction in accuracy. They can be used for initial approximation on massive datasets.
- Cascade SVM gives most relevant support vectors with high accuracy.

# Future work

- Performance of parallel SVM algorithms on unbalanced data
- Quantification of communication cost between mappers and reducers.
- More sophisticated method to calculate training error in Iterative SVM

# Acknowledgement

- Dr. Stan Thomas
- Dr. David John and Dr. William Turkett
- Dr. Todd Torgersen
- Department of Computer Science
- Friends and Family

# Precision comparison

